

AI Ethics Discussion Guide

Facilitated prompts on fairness, bias, and responsible AI



HIGH
SCHOOL

45-60 MIN

DISCUSSION-BASED

Facilitator Notes

This guide is built around four short case studies, each followed by discussion questions with no single 'correct' answer. Encourage students to argue multiple sides. Suggested format: present the case (2 min), small group discussion (5 min), whole-class share-out (5 min), repeat for each case.

Case 1: The Biased Hiring AI

Scenario

A company uses an AI tool to screen job applications. The AI was trained on 10 years of past hiring data. Researchers discover it consistently ranks resumes from certain universities and zip codes lower, even when qualifications are equal.

Discussion questions:

- Why might the AI have learned this pattern, even though no one programmed it to discriminate?
- Who is responsible when an AI causes harm: the company, the programmers, or the data?
- What would you change about how this AI was built or tested before deployment?

Case 2: The Smart Traffic Trade-off

Scenario

A city's AI traffic system is optimized to minimize average wait time citywide. Researchers find this means wealthier neighborhoods (with more cars) get prioritized green lights, while a lower-income neighborhood experiences longer waits more often.

Discussion questions:

- Is 'minimize average wait time' a fair goal for a city-wide system? Why or why not?
- What other goals could the AI be designed to balance instead?
- How would you design a fairer version of this system?

Case 3: The Helpful but Watching Camera

Scenario

Smart streetlights with cameras can detect car accidents and call for help automatically, saving lives. The same cameras can also track which cars pass by and where they go.

Discussion questions:

- How do we weigh the safety benefit against the privacy cost?
- Should communities get to vote on whether these cameras are installed?
- What rules would you put in place about who can access the camera data and for how long?

Case 4: The AI That Won't Explain Itself

Scenario

A bank uses an AI to decide who gets approved for a loan. The AI is very accurate, but even its engineers can't fully explain why it denies any specific application, it's a 'black box.'

Discussion questions:

- Should the bank be allowed to use the AI if it works well but can't be explained?
- What is lost when an important decision can't be explained to the person it affects?
- Where else might 'black box' AI decisions cause real harm if no one can question them?



The Four Pillars

What responsible AI actually requires

Fairness

AI should not produce systematically unfair outcomes for any group of people.

Transparency

People affected by an AI decision should be able to understand, in plain language, how it was made.

Accountability

There must always be a clear human or organization responsible when AI causes harm.

Safety

AI systems, especially in critical infrastructure, must be rigorously tested before deployment.

Closing Reflection (individual, written)

Ask students to write 3-5 sentences responding to: *'Pick one pillar above. Describe one real or imagined AI system where that pillar was missing, and what should have been done differently.'*

